

Introduction to Software Development – CS 6010

Lecture 12 – Number Representations

Master of Software Development (MSD) Program

Varun Shankar

Fall 2023

Brief note on Characters

- A single character is typically stored in 1 Byte (8 bits)
- The computer will “look them up” in the ASCII Table (When displaying)
- Addition works (sort of)
 - $'A' + 3 == 'D'$
- Check to see if a capital letter...
 - if a character is ≤ 90 and ≥ 65 // Note: it is bad to use the ASCII value...
 - Remember, never write 90, or 65. Use 'Z' or 'A'
 - Or 'a' or 'z' or 'm' or '0' (the letter zero) or whatever letter you need.
 - $\text{char} \geq 'A' \ \&\& \ \text{char} \leq 'Z'$
- How to convert the letter '3' to the number 3?
 - $'3' - '0'$
- Convert the letter 'C' to the number 12
 - $'C' - 'A' + 10$

How do we handle signed (-) numbers?

- How do you represent negative numbers in binary?
- One possibility is to represent negative numbers using “One’s Complement”.
- Say for example we have 4 bits to represent both +ve and -ve numbers.
- Now say we want -2.

One's Complement

- Write down 2 in binary using 4 bits: 0010
- Then **invert *all the bits*** : 1101. This is treated as -2.
- The name follows from the fact that if you add a number and its negative, you'll always get 1111 ($0010 + 1101 = 1111$, for example).
 - But this is weird... not 0000! More in a second.
- You can always take the most significant bit and look at it. If it's 1, the number is -ve, if it's 0 the number is +ve.

Weirdness in One's Complement

- With 4 bits (or really any number of bits), when we add a number to its own negative, we get 1111.
- But we'd expect 0000. So what's 1111?
- It's *negative 0*!
- So +0 is different from -0 in one's complement, which is really annoying. Must test for both +0 and -0 if computer is using one's complement.
- Old computers used this (until the 1980s), but most computers use a better scheme called Two's Complement instead.
 - One's complement is now primarily used in older or specialized devices, or completely different contexts than representing negative numbers.

Two's Complement

- Two's complement comes from a really simply idea.
- Say for instance we have 3 bits to represent an integer.
- If we only allow +ve numbers (and 0), this lets us represent the values 0 – 7 (000 = 0, 001 = 1,... 111 = 7).
- If we want to allow both +ve and –ve numbers, one way is to split the range.
 - 000, 001, 010, and 011 are +ve (0, 1, 2, and 3 respectively).
 - 100, 101, 110, and 111 are –ve (-4,-3, -2, and -1 respectively).
 - Notice there's no -0 here.
 - Also notice the +ve numbers all have a 0 in the first bit.

Two's Complement

- In Two's Complement, we specifically find a –ve number using the following algorithm.
- Say we only have 3 bits available and we want to find the binary value of -2. Recall that 2 is 010.
 - First, as in one's complement, flip all the bits: gives 101.
 - Then add 1 to it (and ignore any overflow): gives 110.
 - This is what we saw before on the previous slide!
- What about -0? $000 \rightarrow 111 + 1 \rightarrow 000$ (ignoring overflow).
 - So -0 is the same as 0.

Two's Complement: Bits matter

- The number of bits used to represent the integer really does matter for correctness.
- Say we have 4 bits to represent integers. Then +5 is 0101. What is -5?
 - First, flip all bits. We get 1010.
 - Then add 1. We get 1011.
- But say we have 8 bits instead. +5 is 0000 0101. What is -5?
 - First, flip all bits. We get 1111 1010.
 - Then add 1. We get 1111 1011.
- In all cases, notice negative numbers have 1 for the most significant bit.
 - But... that's weird, right? What about the power expansion?

Two's Complement: Power Expansion??

- Let's go back to 3 bits. The number 2 is 010. -2 is then 110.
- But, $010 = (0 * 2^2) + (1 * 2^1) + (0 * 2^0)$.
- How do we power expand 110? We'll get a wrong answer if we do it naively!
- For two's complement, the power expansion MUST take the sign of the Most significant bit (MSB) into account!
- $110 = (-1 * 2^2) + (1 * 2^1) + (0 * 2^0) = -4 + 2 + 0 = -2$.

Two's Complement: Properties

- MSB acts as both a sign bit and has a value.
- Power expansion takes sign of MSB into account (0 = +ve, 1 = -ve).
- No difference between 0 and -0.
- Number of bits available to represent integer matters!

Two's Complement: More properties

- If you represent an integer with 4 bits, what's the *smallest* number you can represent here?
 - It's 1000, ie, -8. If you replace any of the 0s with 1s, you make the number larger!
 - This is called the **most negative number**.
- What's the **largest** number you can represent?
 - 0111, i.e., 7.
- Sanity check: what do you get if you add these two?
 - $1000 + 0111 = 1111$.
 - What is this number? Be careful, sign bit is 1, so it's -ve.
 - $1111 = (-1 * 2^3) + (1 * 2^2) + (1 * 2^1) + 1 * 2^0 = -8 + 4 + 2 + 1 = -1$.
 - Since $-8 + 7 = -1$, this is indeed correct!

Two's Complement: minus of minus?

- Assume again 3 bits for integers. Recall that -2 is 110.
- What's $-(-2)$? Should be +2!
- Let's check.
 - First, flip all bits. 110 \rightarrow 001.
 - Next, add 1. This gives 010, which is indeed 2!
 - Seems good, right?
- Not quite. There's one edge case: the most negative number.
- Given 3 bits, what is the most negative number?
 - $100 = -1 \cdot 2^2 + 0 \cdot 2^1 + 0 \cdot 2^0 = -4$.
- Let's try and check if $-(-4) = +4$.
 - Flip bits: 100 \rightarrow 011.
 - Add 1: gives 100, but that's just -4 again?!
 - But if we add 100 to 100 and ignore overflow, we get 000...

Two's Complement: Most negative number

- The most negative number **cannot** be negated in two's complement!
- This can result in all sorts of programming errors if you're not careful.
- However, the number added to its negative (which is just itself) still gives 0.

Bit Manipulation

- Tomorrow we will talk about bit manipulation, so it is important to understand what bits represent a number before we start messing with them.

Floating Point Numbers

- What is this?
 - 3.4e12 // It means:
 - $3.4 * 10^{12}$
- This is called Normalized Scientific Notation
 - One non-zero decimal digit before the decimal point.
 - In other words: 34e11 is rewritten as 3.4e12
 - If we had .34e13, we would rewrite it as 3.4e12
- So how are floating point numbers stored in a computer?
 - Basically a binary version of scientific notation
- Significand, Base, Exponent (Using the above 3.4e12)
 - Significand – 3.4
 - Base – In decimal, base is 10 (in Binary the base is 2)
 - Exponent – 12 (Which is used to mean 10^{12})

Floating Point Numbers

- A binary number in scientific notation looks like:
 - $1.01011 * 2^{\text{exp}}$ (where exp is some exponent value)
 - $0.111 * 2^5$
- Since this is *normalized*, what is the only value (in binary) that can come before the decimal point?
 - 1
 - In other words, if we had $0.1111 * 2^{12}$, we would have to write it as?
 - $1.111 * 2^{11}$
 - Because there is always a one in this position, we don't need to store it, and we (ie, the computer) can just assume it is there.
- Making sure we understand this, what does $.01011 * 2^N$ become?
- $1.011 * 2^{N-2}$
 - And technically is stored in the computer as $.011 * 2^{N-2}$

Floating Point Numbers

- $1.011 * 2^N$
- Remember, floating point numbers store the Mantissa (Significand) and the Exponent (the Base is implied).
- Additionally, floating point numbers store a *sign* bit.
- The decimal portion (1.011) is called the Mantissa (Significand)
- Note, there is a sign bit in floating point number storage
- Floats have how many bits?
 - 32 // These 32 bits are allocated like this:
- S Mantissa Exponent (Range: -128 -> 127)



Special Floating Point Numbers

- NaN – Not a Number
 - 0 / 0, etc.
 - Contagious
- +Inf, -Inf
- 0.1 cannot be represented in binary...
 - 0.0001100110011001100110011001100110011001100110011...
- Can't represent many numbers exactly, so we need to know that most floating point numbers are approximated (to some level of accuracy).
- Example in lab today is $.1 + .2 == .3$
 - This turns out to not be true.
 - How can we tell if two floating point number (say A and B) are equal?
 - Just check to see if they are (very) close:
 - $\text{abs}(A - B) < \text{tolerance}$

Unicode (Characters)

- Extension to the ASCII Table (Much Bigger)
 - Includes characters for all written languages, emojis, etc.
- UTF-8
 - ASCII Stuff works the same
 - Use more bytes for higher Unicode characters
 - But we want most text (well standard English characters) to still take only one byte.
 - Compromise:
 - Use 128 values (half the range) with the 1st 128 characters in the ASCII Table.
 - If the character's value is negative, then it is a Unicode character and will use 2-4 total bytes and will have to be decoded.

Midterm Review (i)

- Variables
 - Types, Values, Casting, Declaration, Definition, Assigning, Scope
 - int vs float
 - char * vs string
 - int[] vs vector
 - Manipulating Characters

Midterm Review (ii)

- Asking questions in code
 - if statements, else if, else
 - Boolean logic
- Looping
 - For vs While
 - Syntax
- Strings, Vectors
 - length, substring, indexing into
- Programmer defined datatypes
 - struct
 - Declaring, Creating a variable, accessing fields

Midterm Review (iii)

- Functions
 - Declaration, Definition
 - parameters
 - pass by value
 - pass by reference
 - Returning values from functions
 - return one thing (but it could be a struct and thus hold multiple things)
 - return using parameters (reference params)
 - .h vs .cpp files (multi-file projects / .o files / linking)
- Binary and Hex Numbers
 - Conversion to from decimal numbers

Midterm Review (iv)

- File I/O
 - opening/closing files
 - reading data (strings, numbers, lines)

Starting Practice Problems

- Write a function that takes in a list of the number of cars that passed my house each hour over the last day, and returns the average number of cars that pass my house that day.
- Write a function that takes in 3 integers, subtracts one from each, and returns them - do it in two different ways. [Note, there are 2 ways to “return” values from a function - what are they?]
- Write a function a function that takes in a sentence and returns that sentence without vowels in it. [Note: technically, the returned sentence is a new sentence.]
- Write a ‘for each’ loop that turns all numbers in a list into zero.
- Make sure you are comfortable writing each of the (shorter) functions from the labs and assignments.

Wednesday's Assignment(s)

- Code Review
- Lab – Number Representations