

# Introduction to Software Development – CS 6010

## Lecture 13 – Bitwise Operations

Master of Software Development (MSD) Program

Varun Shankar

Fall 2023

# Miscellaneous

- Midterm Exam is on Friday in class via Gradescope

# Lecture 13 – Bitwise Operations

- Topics
  - Bit Manipulation
  - Boolean Logic
  - Shifting

# Accessing a single Bit

- `Byte == bit[ 8 ].` *// Would like to think of a byte as an array of bits...*
- `Bit firstBit = myByte[ 0 ].` *// Unfortunately C++ does not provide this*
- To manipulate bits, we will use versions of the logical operators
  - `||` and `&&`. *// Boolean operators*
- Specifically, we'll use “bitwise” operators that work on all bits in a variable at the same time.
  - `|` and `&`. *// Note: Only a single | or & when working with bits!*
- Executes a logical Or or And on each bit individually – simultaneously!
- If we are going to work with individual bits, it is almost always a good idea to be specific about the size of variable we are using – i.e.: use types such as `uint8_t` or `int64_t` – preferably an unsigned type.

# Bitwise Boolean Operator Examples

- 0 1 1 0 1 0 1 1      All bits at the
- 0 1 0 1 1 0 1 0 Or ( | ) same time.
- 0 1 1 1 1 0 1 1

- 0 1 1 0 1 0 1 1
- 0 1 0 1 1 0 1 0 And ( & )
- 0 1 0 0 1 0 1 0

- Exclusive Or ( ^ ) One or the other but not both.
- 0 0 | 0
- 0 1 | 1
- 1 0 | 1
- 1 1 | 0

- ▶ 0 1 1 0 1 0 1 1
- ▶ 0 1 0 1 1 0 1 0 Xor ( ^ )
- ▶ 0 0 1 1 0 0 0 1

- ▶ What is the bitwise Unary Operation?
  - ▶ Unary means only has one parameter.
  - ▶ ! Is the Not operator for Booleans
  - ▶ ~ is the Not operator for bits, affects every bit, just like &, |
- ▶ int x = 10101111;
- ▶ x = ~x;

# Bitwise Operations on Integers

- `uint8_t a = 2, b = 4;`
- `uint8_t c = a | b;`
- `010 <- a` // You should start to know the 3 (and even 4) bit patterns
- `100 <- b`
- `110 <- c == 6`
- When we care about what the bits look like, we write them as hex constants:
- `int d = 0xEC; // What is E in bits? C?`
- `1110 <- E`
- `1100 <- C`
- `11101100 <- d`

# Bit Shifting – Move All Bits Left or Right

- Operators << and >>
  - Moves bits to the left or the right (by a given number of bits).
- 0 1 1 0 1 0 1 0 >> 2 // Shift 2 bits to the right – note we fill in with 0s from the left

0 0 0 1 1 0 1 0      Lost...// Most significant bit(s) [left most bit(s)] filled in with 0s\*

- 0 1 1 0 1 0 1 0 << 3

- 0 1 1 0 0 1 0 1 0 0 0 0. // Shifted 3 bits to the left – note we also fill in with 0s from the right

- int a = 0001 1010;
- What does shifting to the left mean? What happens to the number when we shift left?
  - Multiply by 2 (Integer math)
  - But can lose significant bits...
- What does shifting to the right mean?
  - Dividing by 2 (Integer math)

# Bitwise Boolean Identities

- $x \ll 2$ ;
  - Note: This does not change  $x$ !
- $x = x \ll 2$ ;
  - Changes  $x$ . Can also be written as:
  - $x \ll= 2$ ;
- Boolean Identities –  $x$  represents a single bit (which could be a 0 or a 1)
  - $x \mid 1 == 1$
  - $x \mid 0 == x$
  - $x \& 1 == x$
  - $x \& 0 == 0$
- Useful if we want to preserve or change bits. Since all bitwise operations apply to ALL bits in a number simultaneously, we use these identities to be able to keep some bits from changing, and to change others.



# Masking

- `int x = ?;` *// x is any number*
- How could we get just the last 4 bits of x?
- `int low4BitsOfX = x & ?`
- `x & 0x000F` *// Written in hex to make it easier to see...*
- `x & 0000 0000 0000 1111` *// Means this in binary*
- Note `0x000F` is usually written as `0xF`
  
- `x = 1111 1111 1010 0101`
- `0000 0000 0000 1111`
- `0000 0000 0000 0101`


# Two's Complement

- `int8_t = -1;`
- What are the bits?
  - This is a 2's Complement number.
  - 1 1 1 1   1 1 1 1  
      <sup>64</sup>  <sup>16</sup>      <sup>4</sup>  <sup>1</sup>  
      -128  32      8   2
  - What are the values of each bit?
  - But in 2's Complement, the most significant bit has a place value of...
    - -128
  - If the most significant bit is -128, what is the sum of all the other place values?

# Conversion to a bigger type

- `int8_t b = -1; // What does b look like (in bits)?`
- `1 1 1 1 1 1 1 1`
- `int i = b; // what happens to i?`
- How many bits are in i?
  - 32
- Would i become:
- `0000 0000 0000 0000 0000 0000 1111 1111 // This is wrong`
- What is this value? What was b's value?
- `1111 1111 1111 1111 1111 1111 1111 1111 // This is what actually happens.`

# Sign Extension

- When you assign a smaller *signed* integer type into a larger *signed* integer type, the new bits copy the top bit of the smaller value.
- This also applies when you right shift a negative number:
- `1 1 0 0 >> 1` // -4 right shifted == -4 / 2 == -2 [Note 1100 is -8 + 4 => -4]  

- `1 1 1 0` // -8 + 4 + 2 => -2
- `0xF000 >> 4` equals what?
  - `0x0F00` // WAIT! There was no sign extension here! Why?
  - `0xF000` (in a vacuum) is not a negative number. Hex constants are treated as unsigned
  - But if you assign it to a signed number and then shift that int, the signed number will behave as expected
- `short i = 0xF000; // What is the size of a short?`
  - 2 bytes
- `i >> 4 == 0xFF00`
- `0xFF00`

# Bitwise Operations Rule of Thumb

- Rule of Thumb:
  - Use *explicitly typed* unsigned types.
    - So we can ignore sign extension
  - uint8\_t, uint16\_t, uint32\_t, uint64\_t

# Masking

- Think about what bits you wish to change (or read)
- Design a *mask* that represents those numbers. A mask usually contains ones in the bit locations of interest, and 0s in all other locations.

# Let's Do Some Examples

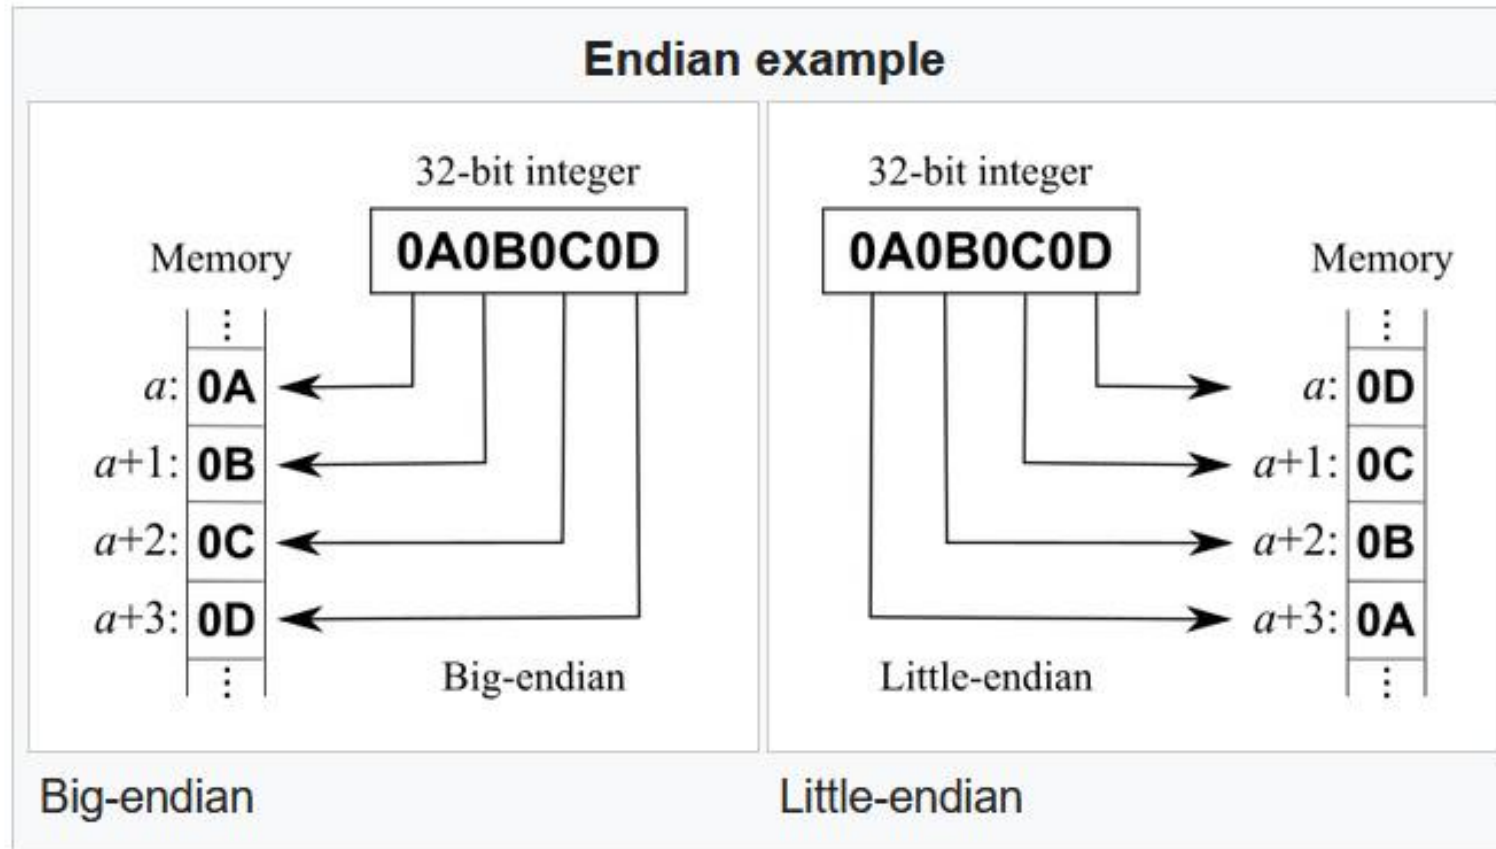
- `void printBits( uint32_t num )`
  - `cout << std::hex << "0x" << num; // To print in Hex`
- `uint32_t x = 0xDEADBEEF;`
- `uint32_t mask = 0xFF000000;`
- How to remove everything but the DE?
  - `x = x & mask;`
  - `x = 0xDE000000`
- How to get just the DE?
  - `x >> 24`
- How to get the BE?
  - `y = x & 0x0000FF00 => y = 0x000BE00`
  - `y = y << 16 => y = 0xBE000000`
  - `Y = y >> 24 => y = 0x000000BE`
  - Could also shift left 16 bits, then shift back right 24 bits.

# Endian-ness

- What date is this:
  - 01/02/03?
  - Is it Jan 2, 2003? Or is it:
  - Feb 1, 2003? Or perhaps:
  - Mar 2, 2001?
  - We have to decide what the ordering means.
- Endian-ness refers to the order of bytes as stored in memory.
- This applies to numbers that are stored in multiple bytes.
  - Which “End” comes first when stored in memory...



# Endianness



# Endianness

- We almost never care...
- So when do we care?
  - Loading binary data from a file (that was created on a different architecture)
  - Reading data off of the network (usually that is coming from a machine that is a different architecture than the one we are using)
- Intel x86, ARMv8 – Little Endian
- PowerPC, older ARM – Big Endian

# Today's Assignment(s)

- Code Review – Book Analyzer
- More catch-up, prep for midterm.
- Bit manipulation assignment tomorrow.